

SUMMARY

ML/AI Engineer with 6+ years turning messy real-world data into production systems that move metrics. Proven in healthcare (~13% insurance denials, LifeMD) and fintech (99.6% anti-fraud efficiency, Sber). Deep expertise in data analysis, NLP, and Generative AI. I build things that work at scale — and stay working.

EXPERIENCE

ML Engineer | LifeMD | Remote | 2023–Present

- Reduced Insurance Denial Rate by 13% by developing a Named Entity Recognition (NER) system based on BioBERT (**AWS Comprehend Medical**) to extract clinical findings from doctor notes.
- Built a patient triage model (XGBoost) with 92% accuracy by integrating structured lab data with features extracted from clinical text.
- Developed automated ETL pipelines (PySpark, Airflow) to process and normalize raw EHR data into a unified format for medical analysis.
- Implemented a Human-in-the-Loop workflow where low-confidence NER predictions were flagged for manual review, ensuring data quality for model retraining.
- Participated in data governance processes: documented data sources, ensured data quality, and maintained data lineage for critical reporting pipelines.

ML Engineer | Sber | Moscow | Remote | 2020–2023

- Reduced False Positive Rate in an **anti-fraud system** by 10% (EFF = **99.6%**) using scikit-learn classifiers, rule induction, and PySpark.
- Built **credit risk models** (XGBoost, Logistic Regression) with **ROC AUC 0.81**. (Links: [1](#), [2](#), [3](#))
- Developed scalable **ETL pipelines** (Spark, Airflow, SQL, Grafana).
- Collaborated with engineering teams on model deployment and **CI/CD**.

Business Analyst & Finance Leadership | Telemarket MD · Cane Factory | Moldova | 2004–2019

Progressed from accountant to Financial Director; led data-driven financial reporting, cost modeling, and business analysis across manufacturing and commercial operations. Built the quantitative foundation for later transition into ML engineering.

PROJECTS

- Linda**: HMM forex scanner + XGBoost Signal Forecaster with live dashboard (Python, DeepSeek, XGBoost, SQLite, Flask, HMM, ATR-based classification, 75 models, 15 symbols × 5 horizons), includes data quality monitoring, automated validation, and business dashboard (Power BI / Flask)
- Fine-tune Turkish ASR Model**: (13.81% WER and 2.96% CER on Turkish Corpus test set)
- ML Project Template**: Production-Ready Machine Learning Platform (FastAPI, Uvicorn, Pandas, Scikit-learn, SQLAlchemy)
- Symptoma**: Real-time symptom matching (**ClinicalBERT**, **UMLS**, **BioBERT**).
- AI HR-Recruiter**: CV-job matching system using **NLP** (**Python**, **Data Scraping**, **TF-IDF**, **Deepseek API**)
- TEST-passer**: **RAG-based assistant** with 88% test accuracy (**Deepseek**, **FAISS**, **RAG**, **Streamlit**).
- Chest X-Ray**: **CNN classifier** for COVID/pneumonia **86%** accuracy (**DS/ML**, **Numpy**, **Keras**, **Sklearn**, **Seaborn**)

EDUCATION

- Stanford University** [CS336: Advanced Language Models](#) | 2024 - audit, in progress

- **Skillbox Deep Learning** | 2024 (ongoing), *Machine Learning Engineer* | 2021–2024
 - **Moscow State Technical University (BMSTU)**, *MSc in Mathematics & Statistics* | 2000–2004
 - **Agrarian University of Moldova**, *BSc in Economics* | 1995–1999
-

SKILLS

- **ML/DL:** Python, Pandas, NumPy, Scikit-learn, PyTorch, TensorFlow, Time Series, CNN, OpenCV, LightAutoML, NLP
- **Generative AI:** LangChain, LlamaIndex, LoRA, RAG, Agents, Agno, n8n
- **Data & MLOps:** PySpark, Airflow, PostgreSQL, FastAPI, Docker, MLflow, Streamlit, CI/CD, model monitoring, scalable pipelines, Data Governance, Data Quality, Data Lineage, Metadata Management, Data Cataloging
- **Tools & Visualization:** Power BI, Matplotlib, Seaborn, FAISS, Git, Grafana, Evidently, AWS, Azure
- **Soft Skills:** Problem-Solving | Technical Communication | Project Leadership
- **Languages:** Russian (native), English (B2), Romanian (B2), Turkish (B2).

Last updated: May, 22, 2026